

Credible consensus among strangers

A jury for funding the research commons

Institute of Open Science Practices (IOSP) 2026, Leiden · Wednesday 14 October 2026, 14:00 to 15:20

Two questions

Did they ship it? Evidence can answer that. How much does it matter? That is a judgment under a policy. Here, credible means public and reproducible; it does not mean true.

The game in five lines

1. Read one published policy.
2. Score each nominee alone; seal the score.
3. Open the scores together.
4. Inside the band, share the reward by closeness to the centre.
5. Outside the band, forfeit more the farther you land from it.

What the design is meant to do

Pay for coherence with the room under one public policy: a Keynesian beauty contest with the judging criteria written down. The curation policy is the Schelling prompt: it asks for honest judgment. Jurors who expect the others to follow it have a reason to follow it too. That shared expectation is the design's bet. Sealed scores keep a visible first score from displacing the prompt. The argument rests on a concrete prompt, jurors who know the field or refuse, and no collusion. Honest jurors can still disagree, and shared bias can still produce agreement. It is a design argument, not a proof.

The band uses the jury's weighted standard deviation; at essentially zero spread, distance forfeits switch off. The workshop uses points and equal seats; draws weighted by stake and appeals belong to the full design.

Workshop simplification: each round with opened scores pays the full base reward, even when those scores all agree.

The shared prompt

How much does this nominee's completed and maintained work underpin how research is done, shared and kept, now?

Docket: Inkscape, Linux, PDF, LaTeX, arXiv, NumPy.

Five objections

Objection	Proposed defence
Herd	Sealed scores hide other jurors' readings until they can no longer change them.
Vote buying	A sealed-score proof can be reported for a bounty before opening, making verification risky.
Collusion	The full design uses commit and reveal, anti-prereveal games and appeals as deterrents. No universal safe threshold is established.
Vague policy	Worked anchors and recusal help strangers apply the same prompt.
Fake identities	In the proposed design, stake weighs seats; concentrated stake remains a concern.

Where it fits

Review panels, grant programmes and community funds: a community sets the policy; a jury supplies scores to inform its decisions.

Try, read, stay in touch



The room



The working paper



Mailing list

Room: `decentralized-curation.0xferit.workers.dev/r/leiden`

Join: `decentralized-curation.0xferit.workers.dev/join/`

Working paper:

`doi.org/10.5281/zenodo.20543760`

Arrow, the first funded round of this design, runs at Octant in early 2027: an invited jury of domain experts scores a real docket under a published policy. Arrow's invited jury is assumed not to collude. `octant.build`